

## predicting occurance of forestfire

```
In [798]: import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score, classification_report, confusion_ma
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
```

loading data

```
In [799]: data = pd.read_csv("C:\\Users\\Niranjan Bhat\\Downloads\\forestfires (1).csv")
data
```

```
Out[799]:
```

	X	Y	month	day	FFMC	DMC	DC	ISI	temp	RH	wind	rain	area
0	7	5	mar	fri	86.2	26.2	94.3	5.1	8.2	51	6.7	0.0	0.00
1	7	4	oct	tue	90.6	35.4	669.1	6.7	18.0	33	0.9	0.0	0.00
2	7	4	oct	sat	90.6	43.7	686.9	6.7	14.6	33	1.3	0.0	0.00
3	8	6	mar	fri	91.7	33.3	77.5	9.0	8.3	97	4.0	0.2	0.00
4	8	6	mar	sun	89.3	51.3	102.2	9.6	11.4	99	1.8	0.0	0.00
...	...	...	...	...	...	...	...	...	...	...	...	...	...
512	4	3	aug	sun	81.6	56.7	665.6	1.9	27.8	32	2.7	0.0	6.44
513	2	4	aug	sun	81.6	56.7	665.6	1.9	21.9	71	5.8	0.0	54.29
514	7	4	aug	sun	81.6	56.7	665.6	1.9	21.2	70	6.7	0.0	11.16
515	1	4	aug	sat	94.4	146.0	614.7	11.3	25.6	42	4.0	0.0	0.00
516	6	3	nov	tue	79.5	3.0	106.7	1.1	11.8	31	4.5	0.0	0.00

517 rows × 13 columns

X - x-axis spatial coordinate within the Montesinho park map: 1 to 9 Y - y-axis spatial coordinate within the Montesinho park map: 2 to 9 month - month of the year: "jan" to "dec" day - day of the week: "mon" to "sun" FFMC - FFMC index from the FWI system: 18.7 to 96.20; Fine fuel moisture code representing the moisture content of litter DMC - DMC index from the FWI system: 1.1 to 291.3; Duff moisture code representing the average moisture content of organic layers and woody material DC - DC index from the FWI system: 7.9 to 860.6; Drought moisture code representing the average moisture content of organic layers ISI - ISI index from the FWI system: 0.0 to 56.10 temp - temperature in Celsius degrees: 2.2 to 33.30 RH - relative humidity in %: 15.0 to 100 wind - wind speed in km/h: 0.40 to 9.40 rain - outside rain in mm/m2 : 0.0 to 6.4 area - the burned area of the forest (in hectares (ha)): 0.00 to 1090.84

## preprocessing

```
In [800]: data.isnull().any()
```

```
Out[800]: X          False
Y          False
month      False
day        False
FFMC       False
DMC        False
DC         False
ISI        False
temp       False
RH         False
wind       False
rain       False
area       False
dtype: bool
```

```
In [801]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517 entries, 0 to 516
Data columns (total 13 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0    X      517 non-null    int64  
 1    Y      517 non-null    int64  
 2   month  517 non-null    object  
 3   day    517 non-null    object  
 4   FFMC   517 non-null    float64 
 5   DMC    517 non-null    float64 
 6   DC     517 non-null    float64 
 7   ISI    517 non-null    float64 
 8   temp   517 non-null    float64 
 9   RH     517 non-null    int64  
10  wind   517 non-null    float64 
11  rain   517 non-null    float64 
12  area   517 non-null    float64 
dtypes: float64(8), int64(3), object(2)
memory usage: 52.6+ KB
```

In [802]: `data.describe()`


Out[802]:

	X	Y	FFMC	DMC	DC	ISI	temp	
count	517.000000	517.000000	517.000000	517.000000	517.000000	517.000000	517.000000	517.0
mean	4.669246	4.299807	90.644681	110.872340	547.940039	9.021663	18.889168	44.2
std	2.313778	1.229900	5.520111	64.046482	248.066192	4.559477	5.806625	16.3
min	1.000000	2.000000	18.700000	1.100000	7.900000	0.000000	2.200000	15.0
25%	3.000000	4.000000	90.200000	68.600000	437.700000	6.500000	15.500000	33.0
50%	4.000000	4.000000	91.600000	108.300000	664.200000	8.400000	19.300000	42.0
75%	7.000000	5.000000	92.900000	142.400000	713.900000	10.800000	22.800000	53.0
max	9.000000	9.000000	96.200000	291.300000	860.600000	56.100000	33.300000	100.0

In [803]: `correlation_matrix = data.corr()`  
`plt.figure(figsize=(12, 8))`  
`sns.heatmap(correlation_matrix, annot=True, cmap='coolwarm', fmt=".2f")`  
`plt.title('Correlation Heatmap')`  
`plt.show()`

C:\Users\Niranjana Bhat\AppData\Local\Temp\ipykernel\_18628\123347052.py:1: FutureWarning: The default value of numeric\_only in DataFrame.corr is deprecated. In a future version, it will default to False. Select only valid columns or specify the value of numeric\_only to silence this warning.

`correlation_matrix = data.corr()`



converting continous values of area column into binary values

```
In [804]: data['fire_occurrence'] = data['area'].apply(lambda x: 1 if x > 0 else 0)
```

```
In [805]: data.fire_occurrence
```

```
Out[805]: 0      0
          1      0
          2      0
          3      0
          4      0
          ..
         512     1
         513     1
         514     1
         515     0
         516     0
          Name: fire_occurrence, Length: 517, dtype: int64
```

converting categorial values into binary

```
In [806]: data = pd.get_dummies(data, columns=['month', 'day', 'rain'])
```

In [807]: data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 517 entries, 0 to 516
Data columns (total 37 columns):
#   Column                Non-Null Count  Dtype
---  -
0   X                      517 non-null   int64
1   Y                      517 non-null   int64
2   FFMC                   517 non-null   float64
3   DMC                    517 non-null   float64
4   DC                     517 non-null   float64
5   ISI                    517 non-null   float64
6   temp                   517 non-null   float64
7   RH                     517 non-null   int64
8   wind                   517 non-null   float64
9   area                   517 non-null   float64
10  fire_occurrence        517 non-null   int64
11  month_apr               517 non-null   uint8
12  month_aug               517 non-null   uint8
13  month_dec               517 non-null   uint8
14  month_feb               517 non-null   uint8
15  month_jan               517 non-null   uint8
16  month_jul               517 non-null   uint8
17  month_jun               517 non-null   uint8
18  month_mar               517 non-null   uint8
19  month_may               517 non-null   uint8
20  month_nov               517 non-null   uint8
21  month_oct               517 non-null   uint8
22  month_sep               517 non-null   uint8
23  day_fri                 517 non-null   uint8
24  day_mon                 517 non-null   uint8
25  day_sat                 517 non-null   uint8
26  day_sun                 517 non-null   uint8
27  day_thu                 517 non-null   uint8
28  day_tue                 517 non-null   uint8
29  day_wed                 517 non-null   uint8
30  rain_0.0                517 non-null   uint8
31  rain_0.2                517 non-null   uint8
32  rain_0.4                517 non-null   uint8
33  rain_0.8                517 non-null   uint8
34  rain_1.0                517 non-null   uint8
35  rain_1.4                517 non-null   uint8
36  rain_6.4                517 non-null   uint8
dtypes: float64(7), int64(4), uint8(26)
memory usage: 57.7 KB
```

In [808]:

```
data.drop('area', axis=1, inplace=True)
```

In [809]:

```
X = data.drop('fire_occurrence', axis=1)
```

In [810]:

```
X.shape
```

Out[810]: (517, 35)

```
In [811]: y = data['fire_occurrence']
```

```
In [812]: y.shape
```

```
Out[812]: (517,)
```

## splitting data

```
In [813]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random
```

```
In [814]: X_train.shape
```

```
Out[814]: (413, 35)
```

```
In [815]: X_test.shape
```

```
Out[815]: (104, 35)
```

```
In [816]: y_train.shape
```

```
Out[816]: (413,)
```

```
In [817]: y_test.shape
```

```
Out[817]: (104,)
```

## Feature scaling

```
In [818]: from sklearn.preprocessing import StandardScaler  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)
```

## Build model

### implementation of Randomforest

```
In [819]: from sklearn.ensemble import RandomForestClassifier  
model = RandomForestClassifier(n_estimators=100, random_state=42)  
model.fit(X_train, y_train)  
probabilities = model.predict_proba(X_test)
```

```
In [820]: #clf = RandomForestClassifier(random_state=42)
```

```
In [821]: model.fit(X_train, y_train)
```

```
Out[821]: RandomForestClassifier(random_state=42)
```

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [822]: y_pred = model.predict(X_test)
```

```
In [823]: accuracy = accuracy_score(y_test, y_pred)
```

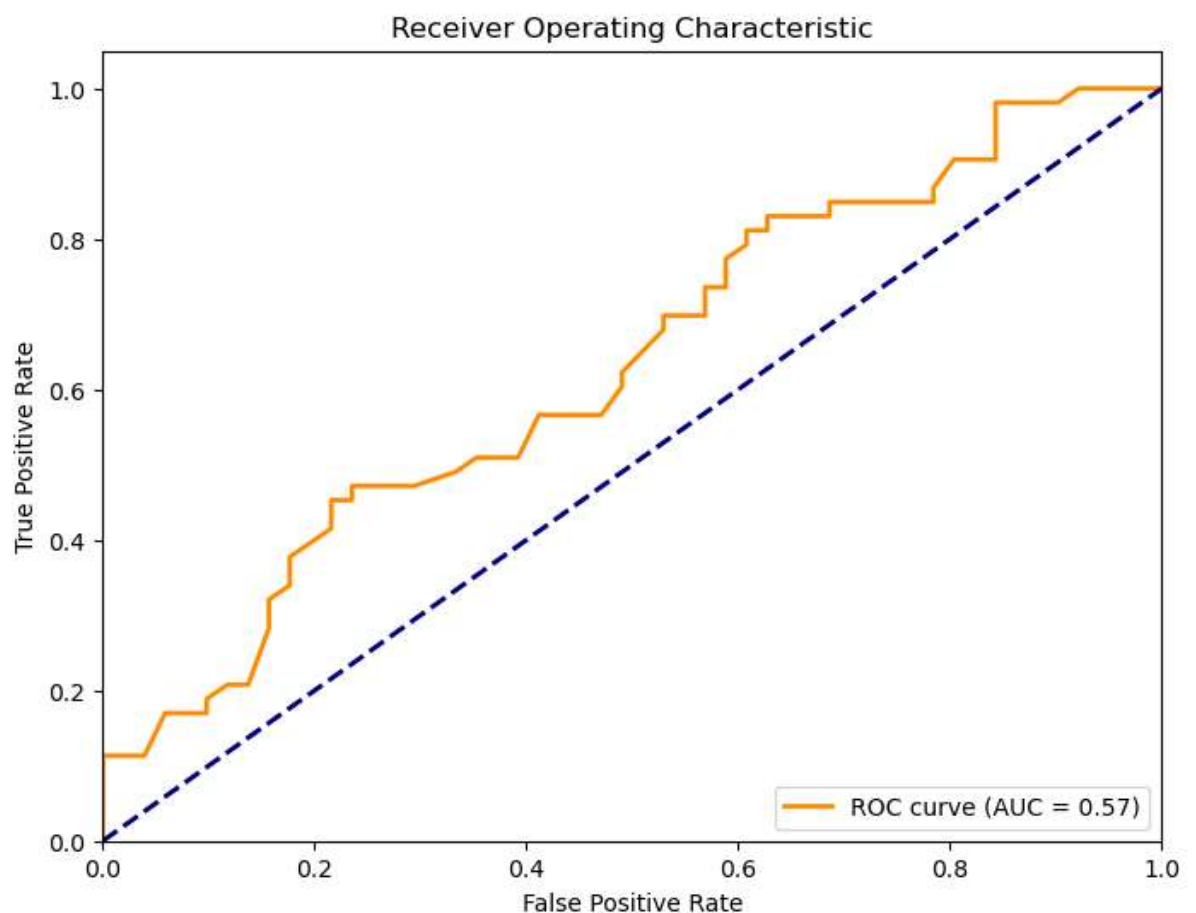
```
In [824]: print("Accuracy:", accuracy)
```

Accuracy: 0.5769230769230769

```
In [825]: from sklearn.metrics import roc_curve, roc_auc_score

fpr, tpr, _ = roc_curve(y_test, model.predict_proba(X_test)[:, 1])
roc_auc = roc_auc_score(y_test, y_pred)

plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (AUC = {roc_auc:0.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc='lower right')
plt.show()
```



```
In [826]: confusion = confusion_matrix(y_test, y_pred)
```

```
In [827]: print("Confusion Matrix:\n", confusion)
```

Confusion Matrix:

```
[[24 27]
 [17 36]]
```



```
In [828]: classification_rep = classification_report(y_test, y_pred)
```

```
In [829]: print("Classification Report:\n", classification_rep)
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.59         0.47         0.52         51
     1       0.57         0.68         0.62         53

 accuracy          0.58         0.57         0.58        104
 macro avg       0.58         0.57         0.57        104
 weighted avg    0.58         0.58         0.57        104
```

## cross validation

```
In [830]: from sklearn.model_selection import cross_val_score, KFold

model = RandomForestClassifier(n_estimators=100, random_state=42)

num_folds = 15

kf = KFold(n_splits=num_folds, shuffle=True, random_state=42)

scores = cross_val_score(model, X, y, cv=kf, scoring='accuracy') # Replace X c

print("Cross-Validation Scores:", scores)

mean_score = scores.mean()
std_score = scores.std()

print(f"Mean Accuracy: {mean_score:.2f}")
print(f"Standard Deviation: {std_score:.2f}")
```

```
Cross-Validation Scores: [0.65714286 0.48571429 0.65714286 0.65714286 0.71428
571 0.6
0.57142857 0.67647059 0.44117647 0.55882353 0.58823529 0.58823529
0.64705882 0.64705882 0.67647059]
Mean Accuracy: 0.61
Standard Deviation: 0.07
```

## implementation of LogisticRegression

```
In [831]: from sklearn.linear_model import LogisticRegression
          clf = LogisticRegression(random_state=42,max_iter=10000)
          clf.fit(X_train, y_train)
```

Out[831]: LogisticRegression(max\_iter=10000, random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

```
In [832]: y_pred = clf.predict(X_test)
```

```
In [833]: accuracy = accuracy_score(y_test, y_pred)
```

```
In [834]: print("Accuracy:", accuracy)
```

Accuracy: 0.5288461538461539

```
In [835]: confusion = confusion_matrix(y_test, y_pred)
```

```
In [836]: print("Confusion Matrix:\n", confusion)
```

Confusion Matrix:

```
[[23 28]
 [21 32]]
```

```
In [837]: classification_rep = classification_report(y_test, y_pred)
          print("Classification Report:\n", classification_rep)
```

Classification Report:

	precision	recall	f1-score	support
0	0.52	0.45	0.48	51
1	0.53	0.60	0.57	53
accuracy			0.53	104
macro avg	0.53	0.53	0.53	104
weighted avg	0.53	0.53	0.53	104

## implementation of DecisionTreeClassifier

In [844]: `from sklearn.tree import DecisionTreeClassifier`

```
# Create a Decision Tree classifier
clf = DecisionTreeClassifier(random_state=42)

# Train the classifier on the training data
clf.fit(X_train, y_train)
```

Out[844]: DecisionTreeClassifier(random\_state=42)

**In a Jupyter environment, please rerun this cell to show the HTML representation or trust the notebook.**

**On GitHub, the HTML representation is unable to render, please try loading this page with nbviewer.org.**

In [845]: `y_pred = clf.predict(X_test)`  
`accuracy = accuracy_score(y_test, y_pred)`  
`print("Accuracy:", accuracy)`

Accuracy: 0.49038461538461536

In [846]: `confusion = confusion_matrix(y_test, y_pred)`  
`print("Confusion Matrix:\n", confusion)`

Confusion Matrix:  
 [[31 20]  
 [33 20]]

In [847]: `classification_rep = classification_report(y_test, y_pred)`  
`print("Classification Report:\n", classification_rep)`

Classification Report:

	precision	recall	f1-score	support
0	0.48	0.61	0.54	51
1	0.50	0.38	0.43	53
accuracy			0.49	104
macro avg	0.49	0.49	0.48	104
weighted avg	0.49	0.49	0.48	104